

```

!pip install haversine

Collecting haversine
  Downloading haversine-2.9.0-py2.py3-none-any.whl.metadata (5.8 kB)
  Downloading haversine-2.9.0-py2.py3-none-any.whl (7.7 kB)
Installing collected packages: haversine
Successfully installed haversine-2.9.0

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import haversine as hs
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score
from sklearn.metrics import mean_squared_error, root_mean_squared_error

df = pd.read_csv('uber.csv')
df

      Unnamed: 0           key  fare_amount \
0        24238194  2015-05-07 19:52:06.0000003    7.5
1        27835199  2009-07-17 20:04:56.0000002    7.7
2        44984355  2009-08-24 21:45:00.00000061   12.9
3        25894730  2009-06-26 08:22:21.0000001    5.3
4        17610152  2014-08-28 17:47:00.000000188   16.0
..          ...
199995     42598914  2012-10-28 10:49:00.00000053    3.0
199996     16382965  2014-03-14 01:09:00.0000008    7.5
199997     27804658  2009-06-29 00:42:00.00000078   30.9
199998     20259894  2015-05-20 14:56:25.0000004   14.5
199999     11951496  2010-05-15 04:08:00.00000076   14.1

      pickup_datetime  pickup_longitude  pickup_latitude \
0  2015-05-07 19:52:06 UTC            -73.999817       40.738354
1  2009-07-17 20:04:56 UTC            -73.994355       40.728225
2  2009-08-24 21:45:00 UTC            -74.005043       40.740770
3  2009-06-26 08:22:21 UTC            -73.976124       40.790844
4  2014-08-28 17:47:00 UTC            -73.925023       40.744085
..          ...
199995  2012-10-28 10:49:00 UTC            -73.987042       40.739367
199996  2014-03-14 01:09:00 UTC            -73.984722       40.736837
199997  2009-06-29 00:42:00 UTC            -73.986017       40.756487
199998  2015-05-20 14:56:25 UTC            -73.997124       40.725452
199999  2010-05-15 04:08:00 UTC            -73.984395       40.720077

      dropoff_longitude  dropoff_latitude  passenger_count

```

```

0           -73.999512    40.723217      1
1           -73.994710    40.750325      1
2           -73.962565    40.772647      1
3           -73.965316    40.803349      3
4           -73.973082    40.761247      5
...
199995       ...          ...          ...
199996       -73.986525    40.740297      1
199997       -74.006672    40.739620      1
199997       -73.858957    40.692588      2
199998       -73.983215    40.695415      1
199999       -73.985508    40.768793      1

```

[200000 rows x 9 columns]

df.head()

	Unnamed: 0	key	fare_amount	\
0	24238194	2015-05-07 19:52:06.0000003	7.5	
1	27835199	2009-07-17 20:04:56.0000002	7.7	
2	44984355	2009-08-24 21:45:00.00000061	12.9	
3	25894730	2009-06-26 08:22:21.0000001	5.3	
4	17610152	2014-08-28 17:47:00.000000188	16.0	
		pickup_datetime	pickup_longitude	pickup_latitude
0	2015-05-07 19:52:06 UTC	-73.999817	40.738354	
1	2009-07-17 20:04:56 UTC	-73.994355	40.728225	
2	2009-08-24 21:45:00 UTC	-74.005043	40.740770	
3	2009-06-26 08:22:21 UTC	-73.976124	40.790844	
4	2014-08-28 17:47:00 UTC	-73.925023	40.744085	
		dropoff_longitude	dropoff_latitude	passenger_count
0	-73.999512	40.723217	1	
1	-73.994710	40.750325	1	
2	-73.962565	40.772647	1	
3	-73.965316	40.803349	3	
4	-73.973082	40.761247	5	

df.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 9 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        200000 non-null  int64  
 1   key               200000 non-null  object  
 2   fare_amount       200000 non-null  float64 
 3   pickup_datetime   200000 non-null  object  
 4   pickup_longitude  200000 non-null  float64 
 5   pickup_latitude   200000 non-null  float64 

```

```

6   dropoff_longitude  199999 non-null  float64
7   dropoff_latitude   199999 non-null  float64
8   passenger_count    200000 non-null  int64
dtypes: float64(5), int64(2), object(2)
memory usage: 13.7+ MB

df.columns

Index(['Unnamed: 0', 'key', 'fare_amount', 'pickup_datetime',
       'pickup_longitude', 'pickup_latitude', 'dropoff_longitude',
       'dropoff_latitude', 'passenger_count'],
      dtype='object')

df = df.drop(['Unnamed: 0', 'key'], axis = 1)

df.shape

(200000, 7)

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 7 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   fare_amount      200000 non-null  float64
 1   pickup_datetime  200000 non-null  object  
 2   pickup_longitude 200000 non-null  float64
 3   pickup_latitude  200000 non-null  float64
 4   dropoff_longitude 199999 non-null  float64
 5   dropoff_latitude  199999 non-null  float64
 6   passenger_count  200000 non-null  int64  
dtypes: float64(5), int64(1), object(1)
memory usage: 10.7+ MB

df.describe()

          fare_amount  pickup_longitude  pickup_latitude
dropoff_longitude \
count  200000.000000      200000.000000      200000.000000
199999.000000
mean      11.359955        -72.527638        39.935885
72.525292
std       9.901776         11.437787        7.720539
13.117408
min     -52.000000       -1340.648410       -74.015515
3356.666300
25%      6.000000         -73.992065        40.734796
73.991407
50%      8.500000         -73.981823        40.752592

```

```
73.980093
75%      12.500000           -73.967154        40.767158      -
73.963658
max      499.000000          57.418457        1644.421482
1153.572603

    dropoff_latitude  passenger_count
count      199999.000000     200000.000000
mean       39.923890        1.684535
std        6.794829        1.385997
min       -881.985513       0.000000
25%        40.733823       1.000000
50%        40.753042       1.000000
75%        40.768001       2.000000
max       872.697628        208.000000

df.isnull().sum()

fare_amount      0
pickup_datetime 0
pickup_longitude 0
pickup_latitude   0
dropoff_longitude 1
dropoff_latitude   1
passenger_count   0
dtype: int64

df.dtypes

fare_amount      float64
pickup_datetime  object
pickup_longitude float64
pickup_latitude   float64
dropoff_longitude float64
dropoff_latitude   float64
passenger_count   int64
dtype: object

df.pickup_datetime = pd.to_datetime(df.pickup_datetime,
                                      errors='coerce')

df.dtypes

fare_amount      float64
pickup_datetime  datetime64[ns, UTC]
pickup_longitude float64
pickup_latitude   float64
dropoff_longitude float64
dropoff_latitude   float64
passenger_count   int64
dtype: object
```

```

df= df.assign(hour = df.pickup_datetime.dt.hour,
              day= df.pickup_datetime.dt.day,
              month = df.pickup_datetime.dt.month,
              year = df.pickup_datetime.dt.year,
              dayofweek = df.pickup_datetime.dt.dayofweek)

df.head()

      fare_amount          pickup_datetime  pickup_longitude
pickup_latitude \
0            7.5 2015-05-07 19:52:06+00:00           -73.999817
40.738354
1            7.7 2009-07-17 20:04:56+00:00           -73.994355
40.728225
2           12.9 2009-08-24 21:45:00+00:00           -74.005043
40.740770
3            5.3 2009-06-26 08:22:21+00:00           -73.976124
40.790844
4           16.0 2014-08-28 17:47:00+00:00           -73.925023
40.744085

      dropoff_longitude  dropoff_latitude  passenger_count  hour  day
month \
0           -73.999512        40.723217             1    19    7
5
1           -73.994710        40.750325             1    20   17
7
2           -73.962565        40.772647             1    21   24
8
3           -73.965316        40.803349             3     8   26
6
4           -73.973082        40.761247             5    17   28
8

      year  dayofweek
0  2015         3
1  2009         4
2  2009         0
3  2009         4
4  2014         3

df = df.drop('pickup_datetime',axis=1)
df.dtypes


fare_amount      float64
pickup_longitude  float64
pickup_latitude   float64
dropoff_longitude  float64
dropoff_latitude   float64
passenger_count      int64

```

```
hour           int32
day            int32
month          int32
year           int32
dayofweek      int32
dtype: object

df.plot(kind = "box", subplots = True, layout = (7,2),
        figsize=(15,20))

fare_amount          Axes(0.125, 0.786098; 0.352273x0.0939024)
pickup_longitude    Axes(0.547727, 0.786098; 0.352273x0.0939024)
pickup_latitude     Axes(0.125, 0.673415; 0.352273x0.0939024)
dropoff_longitude   Axes(0.547727, 0.673415; 0.352273x0.0939024)
dropoff_latitude    Axes(0.125, 0.560732; 0.352273x0.0939024)
passenger_count     Axes(0.547727, 0.560732; 0.352273x0.0939024)
hour               Axes(0.125, 0.448049; 0.352273x0.0939024)
day                Axes(0.547727, 0.448049; 0.352273x0.0939024)
month              Axes(0.125, 0.335366; 0.352273x0.0939024)
year               Axes(0.547727, 0.335366; 0.352273x0.0939024)
dayofweek          Axes(0.125, 0.222683; 0.352273x0.0939024)
dtype: object
```



```

def remove_outlier(df1 , col):
    Q1 = df1[col].quantile(0.25)
    Q3 = df1[col].quantile(0.75)
    IQR = Q3 - Q1
    lower_whisker = Q1-1.5*IQR
    upper_whisker = Q3+1.5*IQR
    df1[col] = np.clip(df1[col] , lower_whisker , upper_whisker)
    return df1

```

```

def treat_outliers_all(df1 , col_list):
    for c in col_list:
        df1 = remove_outlier(df , c)
    return df1

df = treat_outliers_all(df , df.iloc[:, 0::])

travel_dist = []
for pos in range(len(df['pickup_longitude'])):
    long1,lat1,long2,lati2 = [df['pickup_longitude'][pos],
                               df['pickup_latitude'][pos],
                               df['dropoff_longitude'][pos],
                               df['dropoff_latitude'][pos]]
    loc1=(lat1,long1)
    loc2=(lati2,long2)
    c = hs.haversine(loc1,loc2)
    travel_dist.append(c)
print(travel_dist)
df['dist_travel_km'] = travel_dist
df.head()

```

IOPub data rate exceeded.
The Jupyter server will temporarily stop sending output
to the client in order to avoid crashing it.
To change this limit, set the config variable
`--ServerApp.iopub_data_rate_limit`.

Current values:
ServerApp.iopub_data_rate_limit=1000000.0 (bytes/sec)
ServerApp.rate_limit_window=3.0 (secs)

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude \
0	7.5	-73.999817	40.738354	-73.999512
1	7.7	-73.994355	40.728225	-73.994710
2	12.9	-74.005043	40.740770	-73.962565
3	5.3	-73.976124	40.790844	-73.965316
4	16.0	-73.929786	40.744085	-73.973082

	dropoff_latitude	passenger_count	hour	day	month	year
0	40.723217	1.0	19	7	5	2015
3	40.750325	1.0	20	17	7	2009

```

4          40.772647           1.0    21   24     8  2009
2          40.803349           3.0     8   26     6  2009
0
3          40.761247           3.5    17   28     8  2014
4
3

      dist_travel_km
0            1.683325
1            2.457593
2            5.036384
3            1.661686
4            4.116088

df= df.loc[(df.dist_travel_km >= 1) | (df.dist_travel_km <= 130)]
print('Observations left in the dataset:', df.shape)

Observations left in the dataset: (199999, 12)

incorrect_coordinates = df.loc[(df.pickup_latitude > 90) |
                                 (df.pickup_latitude < -90) |
                                 (df.dropoff_latitude > 90) |
                                 (df.dropoff_latitude < -90) |
                                 (df.pickup_longitude > 180) |
                                 (df.pickup_longitude < -180) |
                                 (df.dropoff_longitude > 90) |
                                 (df.dropoff_longitude < -90)]

df.drop(incorrect_coordinates, inplace = True,
       errors = 'ignore')

C:\Users\VEDIKA\AppData\Local\Temp\ipykernel_22680\1102255182.py:1:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
df.drop(incorrect_coordinates, inplace = True,

df.isnull().sum()


fare_amount      0
pickup_longitude 0
pickup_latitude   0
dropoff_longitude 0
dropoff_latitude   0
passenger_count   0
hour              0
day               0

```

```
month          0
year           0
dayofweek      0
dist_travel_km 0
dtype: int64

sns.heatmap(df.isnull())
```

<Axes: >



```
corr = df.corr()
corr
```

	fare_amount	pickup_longitude	pickup_latitude	\
fare_amount	1.000000	0.154056	-0.110856	
pickup_longitude	0.154056	1.000000	0.259492	
pickup_latitude	-0.110856	0.259492	1.000000	
dropoff_longitude	0.218681	0.425622	0.048889	
dropoff_latitude	-0.125874	0.073309	0.515736	


passenger_count	0.015798	-0.013202	-0.012879		
hour	-0.023605	0.011590	0.029691		
day	0.004552	-0.003194	-0.001544		
month	0.030815	0.001168	0.001561		
year	0.141271	0.010193	-0.014247		
dayofweek	0.013664	-0.024645	-0.042304		
dist_travel_km	0.786381	0.048423	-0.073385		
dropoff_longitude dropoff_latitude					
passenger_count \ fare_amount	0.218681	-0.125874			
0.015798					
pickup_longitude	0.425622	0.073309	-		
0.013202					
pickup_latitude	0.048889	0.515736	-		
0.012879					
dropoff_longitude	1.000000	0.245670	-		
0.009304					
dropoff_latitude	0.245670	1.000000	-		
0.006329					
passenger_count	-0.009304	-0.006329			
1.000000					
hour	-0.046560	0.019765			
0.020260					
day	-0.004008	-0.003498			
0.002699					
month	0.002392	-0.001191			
0.010353					
year	0.011347	-0.009595	-		
0.009743					
dayofweek	-0.003337	-0.031932			
0.048542					
dist_travel_km	0.155200	-0.052657			
0.009916					
hour day month year					
dayofweek \ fare_amount	-0.023605	0.004552	0.030815	0.141271	0.013664
pickup_longitude	0.011590	-0.003194	0.001168	0.010193	-0.024645
pickup_latitude	0.029691	-0.001544	0.001561	-0.014247	-0.042304
dropoff_longitude	-0.046560	-0.004008	0.002392	0.011347	-0.003337
dropoff_latitude	0.019765	-0.003498	-0.001191	-0.009595	-0.031932
passenger_count	0.020260	0.002699	0.010353	-0.009743	0.048542
hour	1.000000	0.004664	-0.003924	0.002162	-0.086956

```
day          0.004664  1.000000 -0.017358 -0.012165  0.005609
month        -0.003924 -0.017358  1.000000 -0.115860 -0.008785
year          0.002162 -0.012165 -0.115860  1.000000  0.006116
dayofweek     -0.086956  0.005609 -0.008785  0.006116  1.000000
dist_travel_km -0.035679  0.001738  0.010046  0.022282  0.030403
```

```
fare_amount      dist_travel_km
pickup_longitude 0.786381
pickup_latitude   0.048423
dropoff_longitude -0.073385
dropoff_latitude   0.155200
passenger_count    -0.052657
hour              0.009916
day                -0.035679
month              0.001738
year               0.010046
dayofweek          0.022282
dist_travel_km     0.030403
1.000000
```

```
fig, axis = plt.subplots(figsize = (10,6))
sns.heatmap(df.corr(), annot = True)
```

```
<Axes: >
```



```

x = df[['pickup_longitude', 'pickup_latitude', 'dropoff_longitude',
         'dropoff_latitude', 'passenger_count', 'hour', 'day', 'month',
         'year', 'dayofweek', 'dist_travel_km']]
y = df['fare_amount']

x_train,x_test,y_train,y_test = train_test_split(x,y,test_size = 0.33)
regression = LinearRegression()
regression.fit(x_train,y_train)

LinearRegression()
regression.intercept_
3642.245916151823
regression.coef_
array([ 2.59328347e+01, -7.47376142e+00,   1.94138306e+01, -
       1.76412371e+01,
       6.72633282e-02,   5.33191567e-03,   3.33385157e-03,
      5.89960478e-02,
       3.67958475e-01, -3.39071893e-02,   1.85104129e+00])

```

```
prediction = regression.predict(x_test)
print('Prediction for x:\n', prediction, '\n')
print('Fare Amount test data:\n', y_test)

Prediction for x:
[11.84624141  8.13921591  8.99398439 ...  5.37209863  5.67245379
 8.29924635]

Fare Amount test data:
27708      12.5
6277       10.9
196687      9.7
118237      8.5
194994      6.1
...
39133      11.7
93332      15.0
190275      6.9
169173      4.1
18935      12.0
Name: fare_amount, Length: 66000, dtype: float64

print('R2 Score:\n', r2_score(y_test, prediction))

R2 Score:
0.6616183907981845

MSE = mean_squared_error(y_test, prediction)
print('Mean Squared Error:\n', MSE)

Mean Squared Error:
10.028547305604478

RMSE = root_mean_squared_error(y_test, prediction)
print('Root Mean Squared Error:\n', RMSE)

Root Mean Squared Error:
3.1667881687293953

rf = RandomForestRegressor(n_estimators=100)
rf.fit(x_train, y_train)

y_pred = rf.predict(x_test)
print('Predictions for Fare Amount:\n', y_pred)

Predictions for Fare Amount:
[ 8.782   9.705   8.0947 ... 17.314   6.78    11.53   ] 

R2_Random = r2_score(y_test, y_pred)
print('Random R2 Score:\n', R2_Random)
```

```
Random R2 Score:  
0.7948958964943291
```

```
MSE_Random = mean_squared_error(y_test, y_pred)  
print('Random Mean Squared Error:\n', MSE_Random)
```

```
Random Mean Squared Error:  
6.071112258179303
```

```
RMSE_Random = root_mean_squared_error(y_test, y_pred)  
print('Random Root Mean Squared Error:\n', RMSE_Random)
```

```
Random Root Mean Squared Error:  
2.463962714445838
```